

# Parallel & Distributed Computing: Lecture 9

Alberto Paoluzzi

November 6, 2018

1 Parallel Algorithm Design

2 Memory and Matrices

3 Sources

# Parallel Algorithm Design

# Preliminaries

Source: Introduction to Parallel Computing, Second Edition, By Ananth Grama, Anshul Gupta, George Karypis, Vipin Kumar, Addison Wesley, 2004

Specifying a nontrivial parallel algorithm may include:

- Identifying **portions** of the work that can be **performed concurrently**.
- **Mapping** the concurrent pieces of work onto **multiple processes** running in parallel.
- **Distributing** the input, output, and **the code** associated with the program.
- Managing **accesses to data shared** by multiple processors.
- **Synchronizing** the **processors** at various stages of the **parallel program execution**.

# Decomposition, Tasks, and Dependency Graphs

- Tasks are **programmer-defined units of computation** into which the main computation is subdivided by means of decomposition
- Simultaneous execution of **multiple tasks** is the key to **reducing the time** required to solve the entire problem
- Tasks can be of **arbitrary size**, but once defined, they are regarded as **indivisible units of computation**

## Dependency Graphs

Dependency Graphs abstract the mapping of tasks to processors, and specify where to put synchronization points (barriers)

# Database query processing

## Vehicle database example

ID#	Model	Year	Color	Dealer	Price
4523	Civic	2002	Blue	MN	US\$18,00
3476	Corolla	1999	White	IL	US\$15,00
7623	Camry	2001	Green	NY	US\$21,00
9834	Prius	2001	Green	CA	US\$18,00
6734	Civic	2001	White	OR	US\$17,00
5342	Altima	2001	Green	FL	US\$19,00

# Database query processing

## Vehicle database example

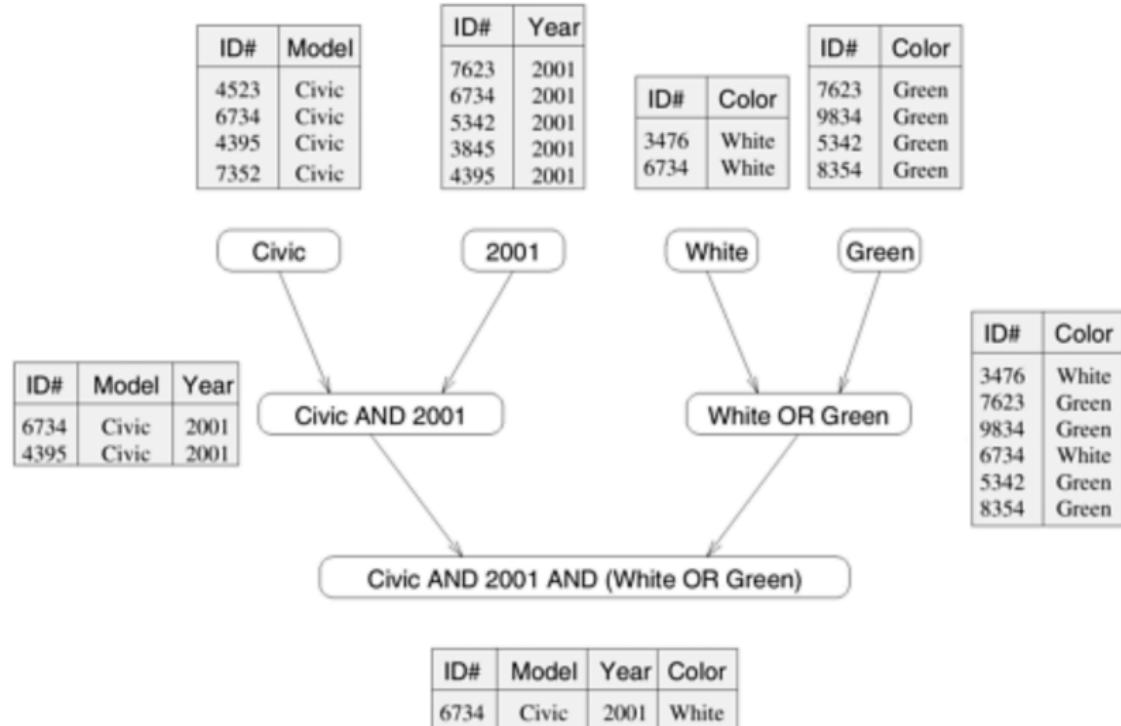
ID#	Model	Year	Color	Dealer	Price
4523	Civic	2002	Blue	MN	US\$18,00
3476	Corolla	1999	White	IL	US\$15,00
7623	Camry	2001	Green	NY	US\$21,00
9834	Prius	2001	Green	CA	US\$18,00
6734	Civic	2001	White	OR	US\$17,00
5342	Altima	2001	Green	FL	US\$19,00

## SQL query example:

MODEL="Civic" AND YEAR="2001" AND (COLOR="Green" OR COLOR="White")

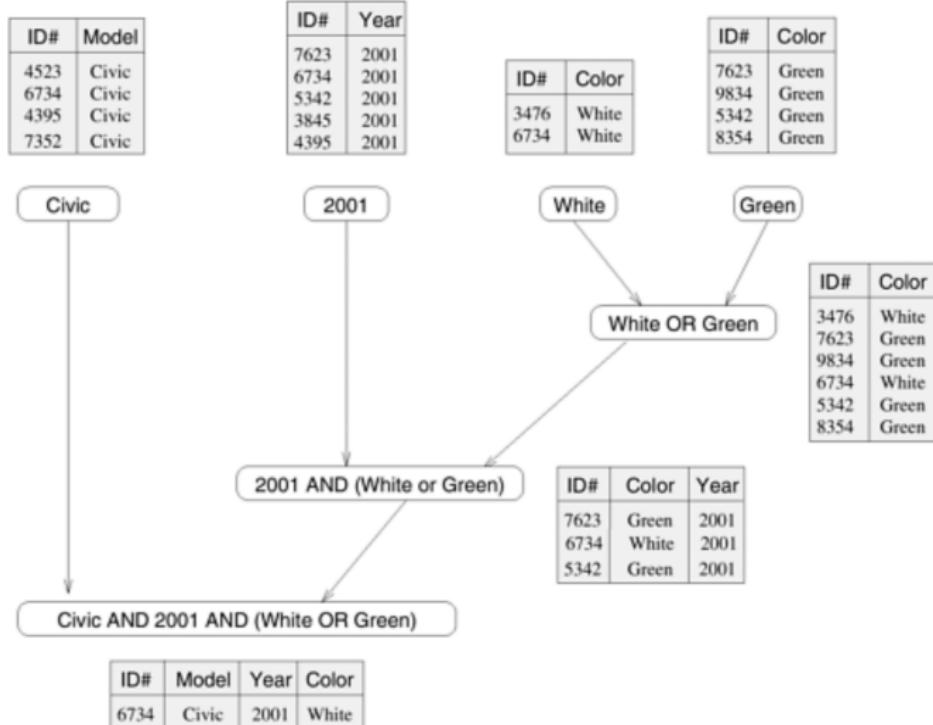
# Database query processing

Data-dependency graph for the query processing operation

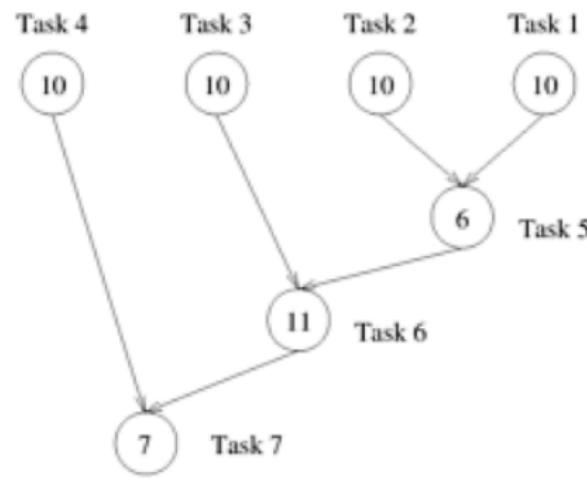
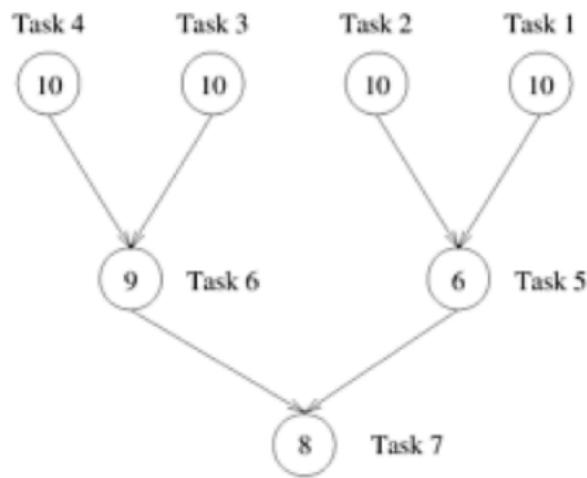


# Database query processing

Another data-dependency graph for the query processing operation



# Abstractions of the task graphs

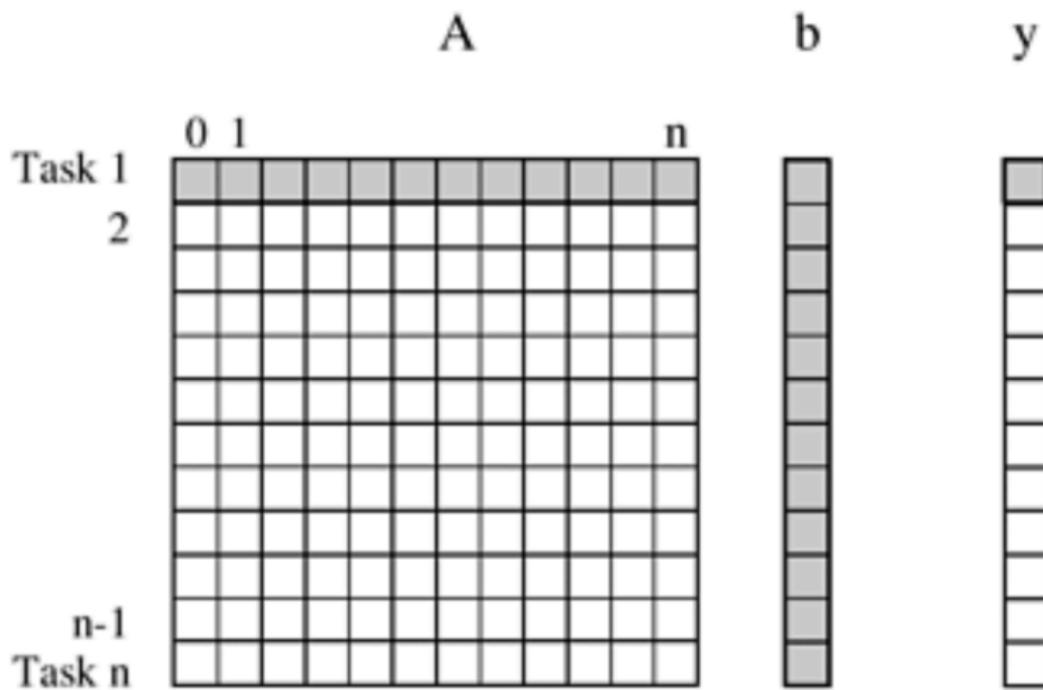


The longest directed path between any pair of start and finish nodes is the **critical path**

The sum of weights of nodes along this path is the **critical path length**

## Dense matrix-vector multiplication

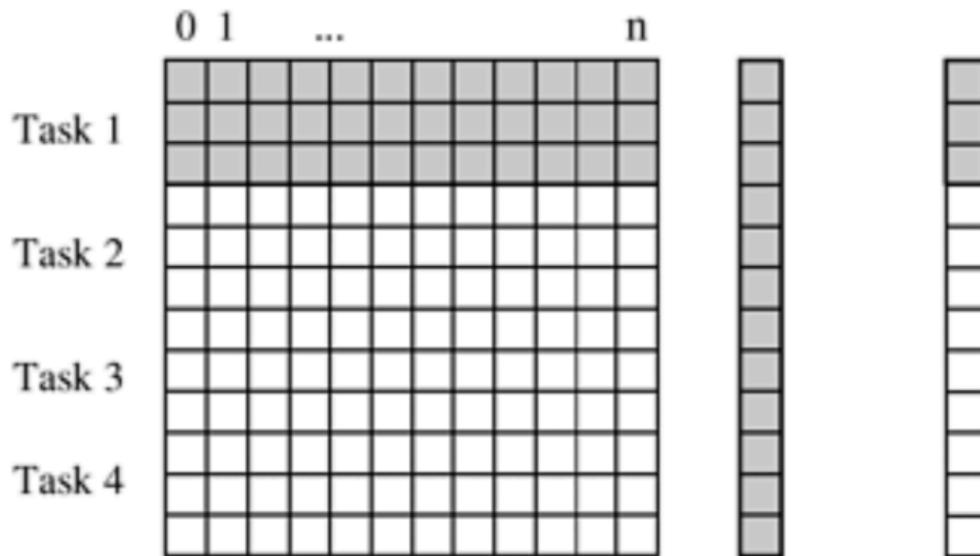
$$\mathbf{A} \mathbf{b} = \mathbf{y}$$



# Granularity

$$\mathbf{A} \mathbf{b} = \mathbf{y}$$

Decomposition of dense matrix-vector multiplication into four tasks



portion of matrix and input/output vectors accessed by Task 1 are highlighted

# Sparse matrix-vector (SpMV) multiplication

Compute  $\mathbf{y} = \mathbf{A}_{n \times n} \times \mathbf{b}$  with  $\mathbf{A}$  sparse matrix

To compute the  $i$ -th entry of the product vector:

$$\mathbf{y}[i] = \sum_{j=1}^n \mathbf{A}[i, j] \mathbf{b}[j]$$

we need to compute the products  $\mathbf{A}[i, j] \times \mathbf{b}[j]$  for only those values of  $j$  for which  $\mathbf{A}[i, j] = 0$

# Sparse matrix-vector (SpMV) multiplication

Compute  $\mathbf{y} = \mathbf{A}_{n \times n} \times \mathbf{b}$  with  $\mathbf{A}$  sparse matrix

To compute the  $i$ -th entry of the product vector:

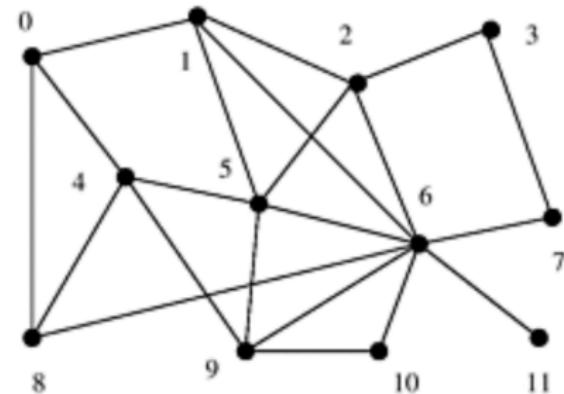
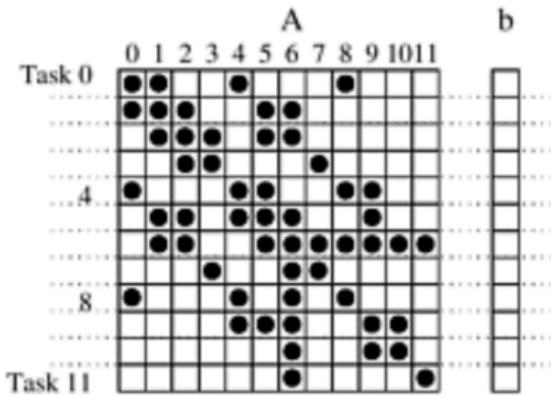
$$\mathbf{y}[i] = \sum_{j=1}^n \mathbf{A}[i, j] \mathbf{b}[j]$$

we need to compute the products  $\mathbf{A}[i, j] \times \mathbf{b}[j]$  for only those values of  $j$  for which  $\mathbf{A}[i, j] \neq 0$

## Protocol (message passing)

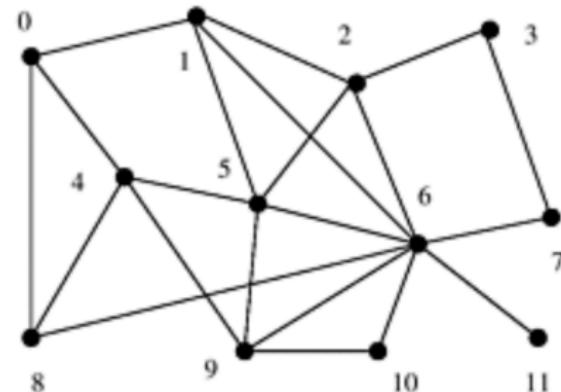
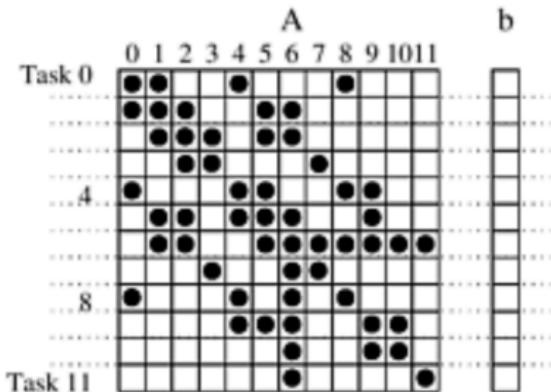
- Task  $i$  computes the element  $y[i]$  of the output vector;
- Task  $i$  is the “owner” of row  $A[i, *]$  and the element  $b[i]$
- Task  $i$  gets the responsibility of sending  $b[i]$  to nodes that need it

# Sparse matrix-vector (SpMV) multiplication



Task 4 must send  $b[4]$  to Tasks 0, 5, 8, and 9 and must get  $b[0]$ ,  $b[5]$ ,  $b[8]$ , and  $b[9]$  to perform its own

# Sparse matrix-vector (SpMV) multiplication



Task 4 must send  $b[4]$  to Tasks 0, 5, 8, and 9 and must get  $b[0]$ ,  $b[5]$ ,  $b[8]$ , and  $b[9]$  to perform its own

## Cohordination tasks

- Initial distribution of input data
- Final collection of output results

# Memory and Matrices

# A trivial problem?

Source: Steven G. Johnson, MIT Applied Math, 18.S096 January 2017

$$C = A \cdot B$$

$$m \times p \quad m \times n \quad n \times p$$

the “obvious” C code (rows · columns):

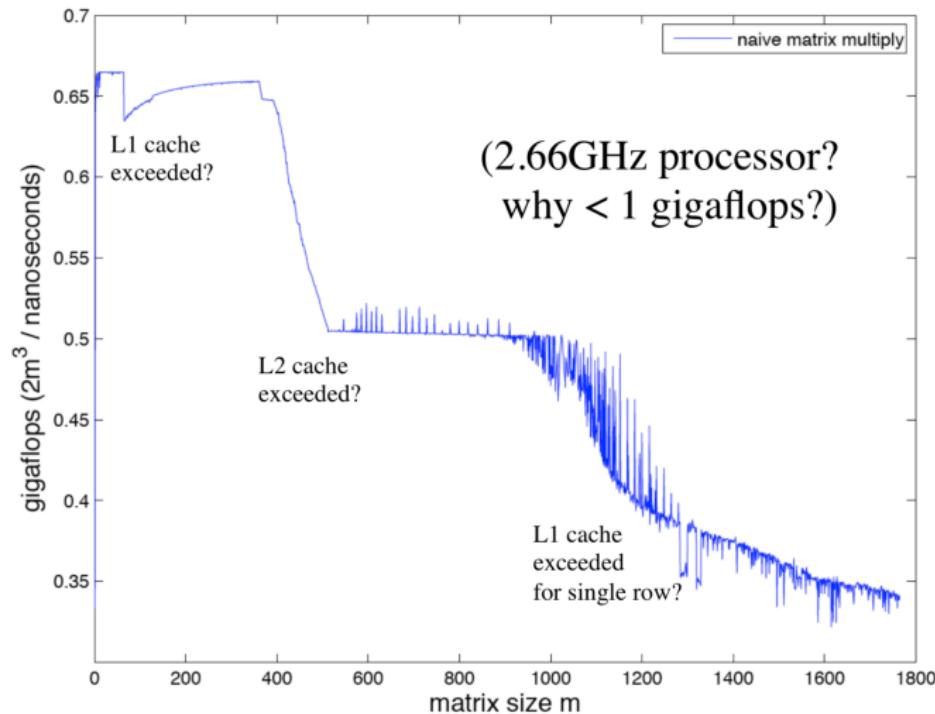
```
/* C = A B, where A is m x n, B is n x p,
   and C is m x p, in row-major order */
void matmul(const double *A, const double *B,
            double *C, int m, int n, int p)
{
    int i, j, k;
    for (i = 0; i < m; ++i)
        for (j = 0; j < p; ++j) {
            double sum = 0;
            for (k = 0; k < n; ++k)
                sum += A[i*n + k] * B[k*p + j];
            C[i*p + j] = sum;
        }
}
```

**for**  $i = 1$  **to**  $m$   
**for**  $j = 1$  **to**  $p$

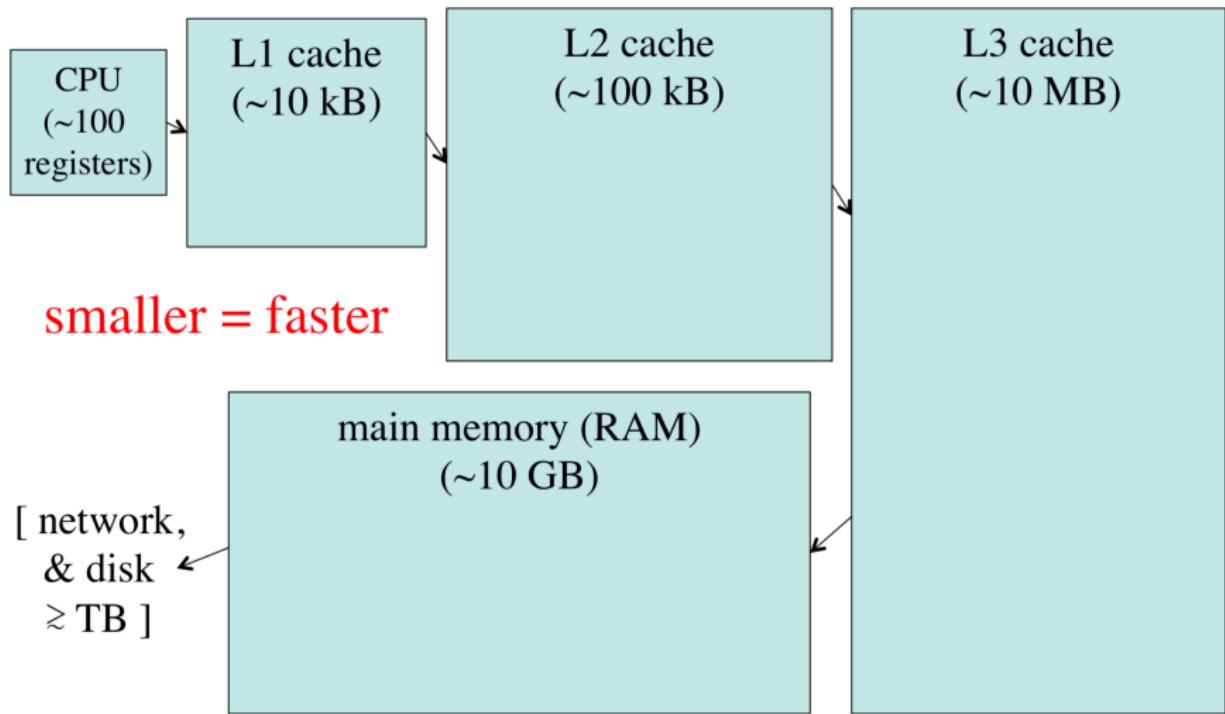
$$C_{ij} = \sum_{k=1}^n A_{ik} B_{kj}$$

**2mnp flops**  
 (adds+mults)  
 “floating-point  
 operations”

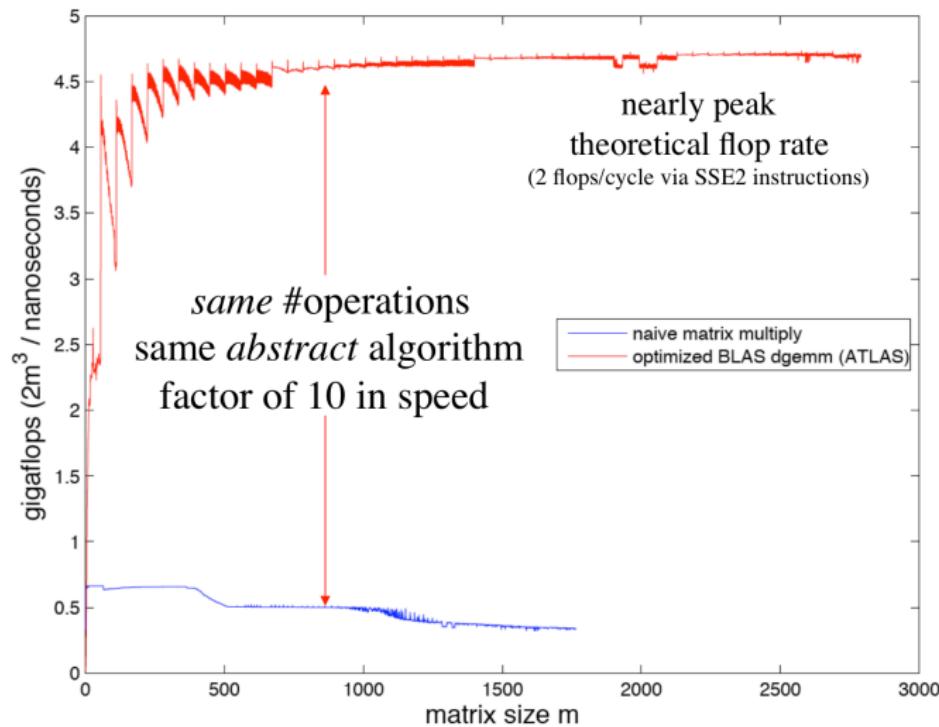
# flops/time is not constant!



# Speed is limited by access to the memory hierarchy



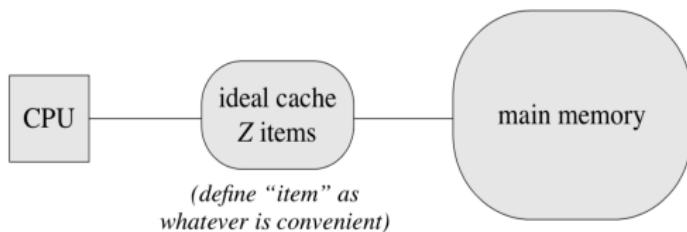
# All flops are not created equal



# Things to remember

- We often **cannot understand performance without understanding memory efficiency** (caches).
  - often ~10 times more important than arithmetic count when working with lots of data
- Computers are **more complicated than you think**.
- Even a trivial algorithm is nontrivial to implement *well*.
  - matrix multiplication: 10 lines of code → **130,000+** (ATLAS)
  - getting the **last factor of 2** in speed often requires wizardry (**and is usually not worth it**)
  - but factors of 10 are often worthwhile and not too hard...

# Ideal Cache Model



when CPU needs an item, either:

- **cache hit:** already in cache (fast)
  - **cache miss:** load into cache (slow)
- goal: analyze # of cache misses

Simplification: **cache is “ideal”**

- **fully associative:** any item in memory can replace any item in cache
- **optimal replacement:** cache miss replaces “best” item  
 (= item not needed for longest time in the future)  
 ... within a ~constant factor of more realistic caches

**Cache complexity:** for problem with  $n$  items, cache size  $Z$ ,  
 want # misses for large  $n$  in “big O” or “big Θ” notation  
 (ignoring constant factors), e.g.  $\Theta(n/Z)$

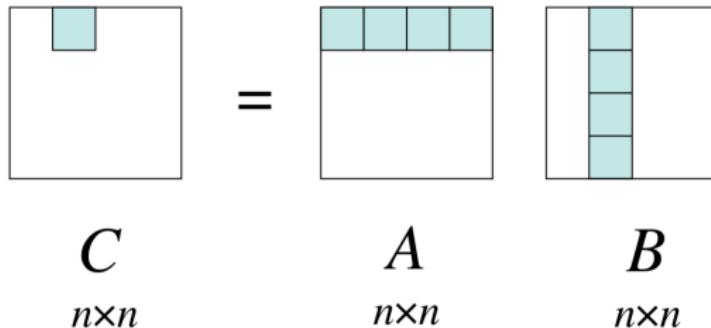
# Strategy for efficient cache utilization: Maximize temporal locality

*Strategy for efficient cache utilization:*  
**Maximize temporal locality**

Once we read an item from memory, we want as much computation as possible before reading the next item...

Re-arrange our algorithm so that computations on the same data occur at close to the same time.

# (optimal) Blocked Matrix Multiply



divide matrices into  $b \times b$  blocks of  $b = \sqrt{Z}/3$  numbers  
 load 3  $b \times b$  blocks into cache and *multiply blocks in-cache*

(provably optimal) cache misses:  $\Theta(n^3/b^3) \times \Theta(b^2) = \Theta(n^3/\sqrt{Z})$

$\# \text{ block} \times \text{block}$	$\# \text{ cache misses}$
$\text{multiplications}$	$\text{per block}$

# (optimal) Cache-Oblivious Matrix Multiply

$$\begin{pmatrix} \mathbf{C}_{1,1} & \mathbf{C}_{1,2} \\ \mathbf{C}_{2,1} & \mathbf{C}_{2,2} \end{pmatrix} = \begin{pmatrix} \mathbf{0}_{1,1} & \mathbf{0}_{1,2} \\ \mathbf{0}_{2,1} & \mathbf{0}_{2,2} \end{pmatrix} + \begin{pmatrix} \mathbf{A}_{1,1} & \mathbf{A}_{1,2} \\ \mathbf{A}_{2,1} & \mathbf{A}_{2,2} \end{pmatrix} \times \begin{pmatrix} \mathbf{B}_{1,1} & \mathbf{B}_{1,2} \\ \mathbf{B}_{2,1} & \mathbf{B}_{2,2} \end{pmatrix}$$

# (optimal) Cache-Oblivious Matrix Multiply

$$\begin{pmatrix} \mathbf{C}_{1,1} & \mathbf{C}_{1,2} \\ \mathbf{C}_{2,1} & \mathbf{C}_{2,2} \end{pmatrix} = \begin{pmatrix} \mathbf{0}_{1,1} & \mathbf{0}_{1,2} \\ \mathbf{0}_{2,1} & \mathbf{0}_{2,2} \end{pmatrix} + \begin{pmatrix} \mathbf{A}_{1,1} & \mathbf{A}_{1,2} \\ \mathbf{A}_{2,1} & \mathbf{A}_{2,2} \end{pmatrix} \times \begin{pmatrix} \mathbf{B}_{1,1} & \mathbf{B}_{1,2} \\ \mathbf{B}_{2,1} & \mathbf{B}_{2,2} \end{pmatrix}$$

$$\mathbf{C}_{1,1} = \mathbf{0}_{1,1} + \mathbf{A}_{1,1} \times \mathbf{B}_{1,1} + \mathbf{A}_{1,2} \times \mathbf{B}_{2,1}$$

$$\mathbf{C}_{1,2} = \mathbf{0}_{1,2} + \mathbf{A}_{1,1} \times \mathbf{B}_{1,2} + \mathbf{A}_{1,2} \times \mathbf{B}_{2,2}$$

$$\mathbf{C}_{2,1} = \mathbf{0}_{2,1} + \mathbf{A}_{2,1} \times \mathbf{B}_{1,1} + \mathbf{A}_{2,2} \times \mathbf{B}_{2,1}$$

$$\mathbf{C}_{2,2} = \mathbf{0}_{2,2} + \mathbf{A}_{2,1} \times \mathbf{B}_{1,2} + \mathbf{A}_{2,2} \times \mathbf{B}_{2,2}$$

# (optimal) Cache-Oblivious Matrix Multiply

$$\begin{pmatrix} \mathbf{C}_{1,1} & \mathbf{C}_{1,2} \\ \mathbf{C}_{2,1} & \mathbf{C}_{2,2} \end{pmatrix} = \begin{pmatrix} \mathbf{0}_{1,1} & \mathbf{0}_{1,2} \\ \mathbf{0}_{2,1} & \mathbf{0}_{2,2} \end{pmatrix} + \begin{pmatrix} \mathbf{A}_{1,1} & \mathbf{A}_{1,2} \\ \mathbf{A}_{2,1} & \mathbf{A}_{2,2} \end{pmatrix} \times \begin{pmatrix} \mathbf{B}_{1,1} & \mathbf{B}_{1,2} \\ \mathbf{B}_{2,1} & \mathbf{B}_{2,2} \end{pmatrix}$$

$$\mathbf{C}_{1,1} = \mathbf{0}_{1,1} + \mathbf{A}_{1,1} \times \mathbf{B}_{1,1} + \mathbf{A}_{1,2} \times \mathbf{B}_{2,1}$$

$$\mathbf{C}_{1,2} = \mathbf{0}_{1,2} + \mathbf{A}_{1,1} \times \mathbf{B}_{1,2} + \mathbf{A}_{1,2} \times \mathbf{B}_{2,2}$$

$$\mathbf{C}_{2,1} = \mathbf{0}_{2,1} + \mathbf{A}_{2,1} \times \mathbf{B}_{1,1} + \mathbf{A}_{2,2} \times \mathbf{B}_{2,1}$$

$$\mathbf{C}_{2,2} = \mathbf{0}_{2,2} + \mathbf{A}_{2,1} \times \mathbf{B}_{1,2} + \mathbf{A}_{2,2} \times \mathbf{B}_{2,2}$$

8 block-sums & 8 block-products

# (optimal) Cache-Oblivious Matrix Multiply

$$\begin{array}{ccc}
 \begin{array}{|c|c|c|c|} \hline
 & & & \\ \hline
 \end{array} & = & 
 \begin{array}{|c|c|c|c|} \hline
 & & & \\ \hline
 \end{array} \\
 C & A & B \\
 n \times n & n \times n & n \times n
 \end{array}$$

*divide and conquer:*

divide  $C$  into 4 blocks

compute block multiply recursively

achieves optimal  $\Theta(n^3/\sqrt{Z})$  cache complexity  
 without knowing the  $Z$ , works for nested caches too

# (optimal) Cache-Oblivious Matrix Multiply

Source: Steven G. Johnson, MIT Applied Math, 18.S096 January 2017

```

function add_matmul_rec!(m,n,p, i0,j0,k0, C,A,B)
    if m+n+p <= 64    # base case: naive matmult for sufficiently large matrices
        for i = 1:m
            for k = 1:p
                c = zero(eltype(C))
                for j = 1:n
                    @inbounds c += A[i0+i,j0+j] * B[j0+j,k0+k]
                end
                @inbounds C[i0+i,k0+k] += c
            end
        end
    else
        m2 = m + 2; n2 = n + 2; p2 = p + 2
        add_matmul_rec!(m2, n2, p2, i0, j0, k0, C, A, B)

        add_matmul_rec!(m-m2, n2, p2, i0+m2, j0, k0, C, A, B)
        add_matmul_rec!(m2, n-n2, p2, i0, j0+n2, k0, C, A, B)
        add_matmul_rec!(m2, n2, p-p2, i0, j0, k0+p2, C, A, B)

        add_matmul_rec!(m-m2, n-n2, p2, i0+m2, j0+n2, k0, C, A, B)
        add_matmul_rec!(m2, n-n2, p-p2, i0, j0+n2, k0+p2, C, A, B)
        add_matmul_rec!(m-m2, n2, p-p2, i0+m2, j0, k0+p2, C, A, B)

        add_matmul_rec!(m-m2, n-n2, p-p2, i0+m2, j0+n2, k0+p2, C, A, B)
    end
    return C
end

```

# Sources

# Sources of lecture materials

- ① Grama, Gupta, Karypis, Kumar [Introduction to Parallel Computing](#)
- ② 18.S096, IAP 2018: [Performance Computing in a High Level Language](#)
- ③ Steven G. Johnson, [More Dots: Syntactic Loop Fusion in Julia](#)