

# Parallel & Distributed Computing: Lecture 19

Alberto Paoluzzi

December 7, 2017



# Parallel computing: Channels

Examples from

[\*<https://github.com/cvdlab/LARLIB.jl>\*](<https://github.com/cvdlab/LARLIB.jl>)

- 1 Algorithm: spatial arrangement from geometric data
- 2 Produce & Consume paradigm: Channels in Julia
- 3 Parallel implementation: Spatial\_arrangement in LARLIB.jl

# Algorithm: spatial arrangement from geometric data

# Arrangement

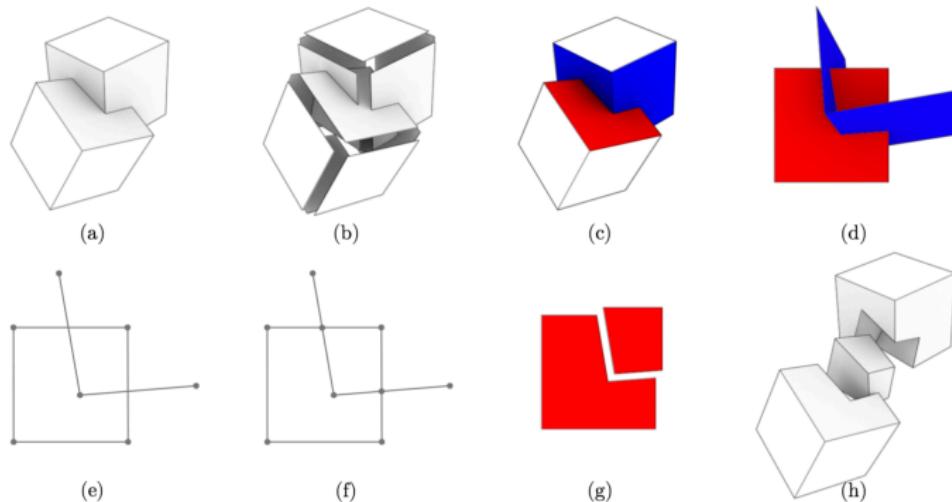


Fig. 1. Cartoon display of the Merge algorithm: (a) the two input solids; (b) the exploded input collection  $B$  of 2-cells embedded in  $\mathbb{E}^3$ ; (c) 2-cell  $\sigma$  (red) and the set  $\Sigma(\sigma)$  (blue) of possible intersection; (d) affine map of  $\sigma \cup \Sigma$  on  $z = 0$  plane; (e) reduction to a set of 1D segments in  $\mathbb{E}^2$ ; (f) pairwise intersections; (g) regularized plane arrangement  $\mathcal{A}(\sigma \cup \Sigma)$ ; (h) exploded 3-complex extracted from the 2-complex  $X_2(\cup_{\sigma \in B} \mathcal{A}(\sigma \cup \Sigma))$  in  $\mathbb{E}^3$ . The LAR representation of both the input data (not a complex) and the output data (3-complex with three 3-cells) is given in Appendix A

Figure 1: Arrangement of  $\mathbb{E}^3$  generated by a collection of 2-complexes

# Arrangement

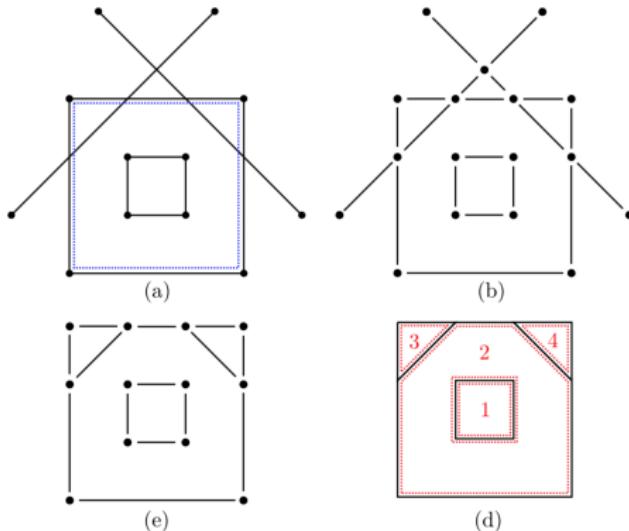


Fig. 2. Basic case: computation of the regularized arrangement of a set of lines in  $\mathbb{E}^2$ : (a) the input, i.e. the 2-cell  $\sigma$  (blue) and the line segment intersections of  $\Sigma(\sigma)$  with  $z = 0$ ; (b) all pairwise intersections; (c) removal of the 1-subcomplex external to  $\sigma$ . To this purpose an efficient and robust point classification algorithm [?] is used, ; (d) interior cells of the regularized 2-complex  $X_2 = \mathcal{A}(\sigma \cup \Sigma)$  generated as arrangement of  $\mathbb{E}^2$  produced by  $\sigma \cup \Sigma$ .

# Arrangement

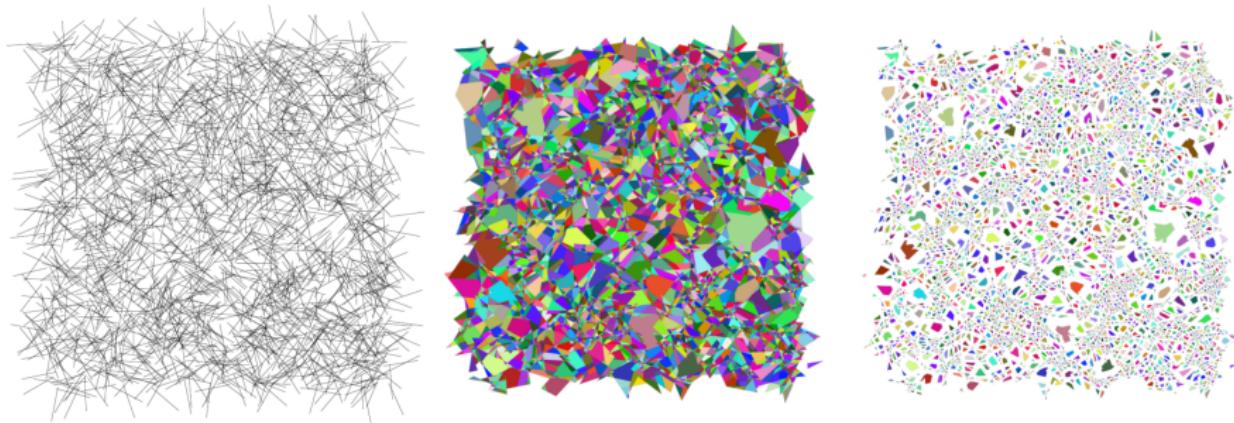


Fig. 5. The regularized 2D arrangement  $X_2$  of the plane generated by a set of random line segments. Note that cells  $\sigma \in X_2$  are not necessarily convex. The Euler characteristic is  $\chi = \chi_0 - \chi_1 + \chi_2 = 11361 - 20813 + 9454 = 2$ .

Figure 3: Arrangement of 2D plane generated by random line segments

# Arrangement

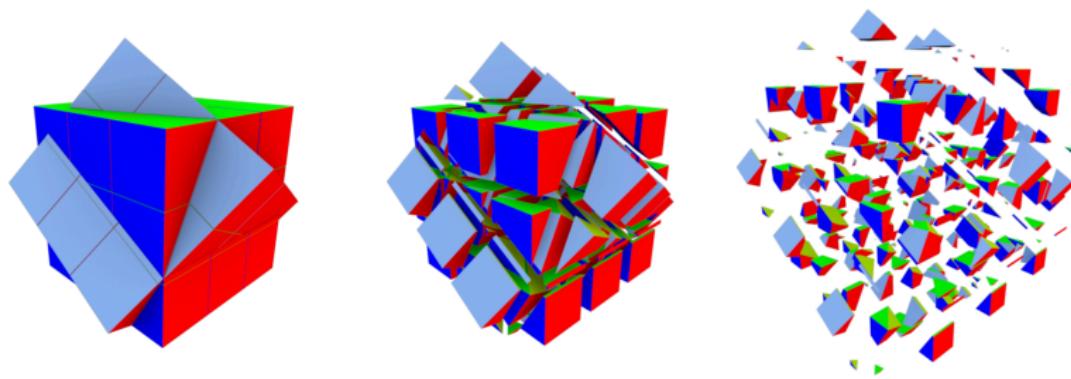


Fig. 6. Merge of two rotated  $3 \times 3 \times 3$  cuboidal 3-complexes: (a) the initial positions of the two input 3-complexes, each with  $3^3$  unit cubic cells; (b) 3-cells of merged complex, with space explosion of small scaling parameter; (c) larger space explosion. Each exploded cell is translated by a scaling-induced movement of its centroid. Output cells are not necessarily convex. In the merged complex we get 236 three-cells and 816 two-cells, both possibly non-convex. Ratio new/old 3-cells is 4.3704.

Figure 4: Arrangement of 3 space generated by two Rubik cubes (2-complexes)

# Arrangement

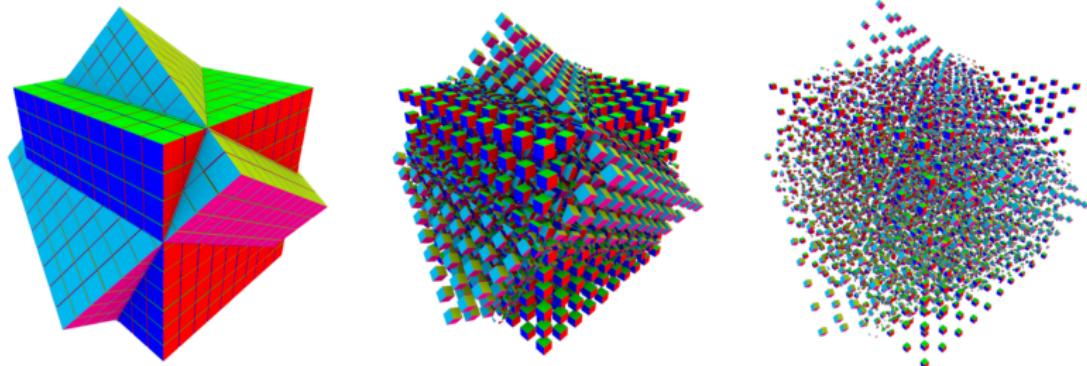


Fig. 7. Merge of two  $10^3$  3-complexes. In the merged complex we get 8655 3-cells, 26600 2-cells (faces), 26732 1-cells (edges), and 8787 0-cells (vertices). Please note that the Euler characteristic is  $\chi = \chi_0 - \chi_1 + \chi_2 - \chi_3 = 8787 - 26732 + 26600 - 8655 = 0$ , since the decomposed 3-complex (no its exploded image!) is homeomorphic to the  $n$ -sphere, where  $\chi = 1 + (-1)^n$ , with  $n = 3$ . Ratio new/old 3-cells is 4.3275.

Figure 5: Arrangement of 3 space generated by two  $10 \times 10 \times 10$  cubes (2-complexes)

# Algorithm: fragmentation of 2-cells

---

**ALGORITHM 3:** Fragmentation of 2-cells
 

---

**Input:**  $\mathcal{S}_2 \subset \mathcal{S}_{d-1}$  # collection of all 2-cells from  $\mathcal{S}_{d-1}$  input in  $\mathbb{E}^d$

**Output:**  $[\partial_2]$  # CSC signed matrix

$\widetilde{\mathcal{S}}_2 = \emptyset$  # initialisation of collection of local fragments

**for**  $\sigma \in \mathcal{S}_2$  **do** # for each 2-cell  $\sigma$  in the input set

- $M = SubManifoldMap(\sigma)$  # affine transform s.t.  $\sigma \mapsto x_3 = 0$  subspace
- $\Sigma = M\mathcal{I}(\sigma)$  # apply the transformation to (possible) incidencies to  $\sigma$
- $\mathcal{S}_1(\sigma) = \emptyset$  # collection of line segments in  $x_3 = 0$
- for**  $\tau \in \Sigma$  **do** # for each 2-cell  $\tau$  in  $\Sigma$ 
  - $\mathcal{P}(\tau), \mathcal{L}(\tau) = \emptyset, \emptyset$  # intersection points and int. segment(s) with  $x_3 = 0$
  - for**  $\lambda \in X_1(\tau)$  **do** # for each 1-cell  $\lambda$  in  $X_1(\tau)$ 
    - if**  $\lambda \notin \{q \mid x_3(q) = 0\}$  **then**  $\mathcal{P}(\tau) += \{p\}$  # append the intersection point of  $\lambda$  with  $x_3 = 0$
- end**
- $\mathcal{L}(\tau) = Points2Segments(\mathcal{P}(\tau))$  # Compute a set of collinear intersection segments
- $\mathcal{S}_1(\sigma) += \mathcal{L}(\tau)$  # accumulate intersection segments with  $\sigma$  generated by  $\tau$

- end**
- $X_2(\sigma) = \mathcal{A}(\mathcal{S}_1(\sigma))$  # arrangement of  $\sigma$  space induced by a soup of 1-complexes
- $\widetilde{\mathcal{S}}_2 += M^{-1}X_2$  # accumulate local fragments, back transformed in  $\mathbb{E}^d$

**end**

$[\partial_1] = QuotientBases(\widetilde{\mathcal{S}}_2)$  # identification of 0- and 1-cells using *kd*-trees and canonical LAR

$[\partial_2] = Algorithm\_I([\partial_1])$  # output computation GRANDE STRONZATA

**return**  $[\partial_2]$

---

# Algorithm: Reconstruction of $\partial_3$ matrix

---

**ALGORITHM 1:** Computation of signed  $[\partial_d^+]$  matrix
 

---

```

/* Pre-condition: d equal to space dimension, s.t.  $(d - 1)$ -cells are shared by two  $d$ -cells */  

/*  

Input:  $[\partial_{d-1}]$  # CSC signed matrix  

Output:  $[\partial_d^+]$  # CSC signed matrix  

 $m, n = [\partial_{d-1}].size; marks = Zeros(n)$  # initializations  

while Sum(marks) <  $2n$  do  

     $\sigma = Choose(marks)$  # select the  $(d - 1)$ -cell seed of the column extraction  

    if marks[ $\sigma$ ] == 0 then  $[c_{d-1}] = [\sigma]$   

    else if marks[ $\sigma$ ] == 1 then  $[c_{d-1}] = [-\sigma]$   

     $[c_{d-2}] = [\partial_{d-1}] [c_{d-1}]$  # compute boundary  $c_{d-2}$  of seed cell  

    while  $[c_{d-2}] \neq []$  do # loop until boundary becomes empty  

        corolla = []  

        for  $\tau \in c_{d-2}$  do # for each hinge  $\tau$  cell  

             $[b_{d-1}] = [\tau]^t [\partial_{d-1}]$  # compute the  $\tau$  coboundary  

            pivot =  $\{|b_{d-1}|\} \cap \{|c_{d-1}|\}$  # compute the  $\tau$  support  

            if  $\tau > 0$  then adj = Next(pivot, Ord( $b_{d-1}$ )) # compute the new adj cell  

            else if  $\tau < 0$  then adj = Prev(pivot, Ord( $b_{d-1}$ ))  

            if  $\partial_{d-1}[\tau, adj] \neq \partial_{d-1}[\tau, pivot]$  then corolla[adj] =  $c_{d-1}[pivot]$  # orient adj  

            else corolla[adj] =  $-(c_{d-1}[pivot])$   

        end  

         $[c_{d-1}] += corolla$  # insert corolla cells in current  $c_{d-1}$   

         $[c_{d-2}] = [\partial_{d-1}] [c_{d-1}]$  # compute again the boundary of  $c_{d-1}$   

    end  

    for  $\sigma \in c_{d-1}$  do marks[ $\sigma$ ] += 1 # update the counters of used cells  

     $[\partial_d^+] += [c_{d-1}]$  # append a new column to  $[\partial_d^+]$   

end  

return  $[\partial_d^+]$ 
  
```

---

# Produce & Consume paradigm: Channels in Julia

# Definition: produce and consume problem

- Also known as the **bounded-buffer problem** is a classic example of a multi-process synchronization problem.

# Definition: produce and consume problem

- Also known as the **bounded-buffer problem** is a classic example of a multi-process synchronization problem.
- Two processes, the **producer** and the **consumer**, who share a common, **fixed-size buffer** used as a **queue** (FIFO)

# Definition: produce and consume problem

- Also known as the **bounded-buffer problem** is a classic example of a multi-process synchronization problem.
- Two processes, the **producer** and the **consumer**, who share a common, **fixed-size buffer** used as a **queue** (FIFO)
- The **producer's job** is to **generate data**, put it into the buffer, and start again.

# Definition: produce and consume problem

- Also known as the **bounded-buffer problem** is a classic example of a multi-process synchronization problem.
- Two processes, the **producer** and the **consumer**, who share a common, **fixed-size buffer** used as a **queue** (FIFO)
- The **producer's job** is to **generate data**, put it into the buffer, and start again.
- At the same time, the **consumer is consuming the data** (i.e., removing it from the buffer), one piece at a time.

## Solution: produce and consume problem

- The solution for the **producer** is to either **go to sleep** or discard data if the buffer is full.

# Solution: produce and consume problem

- The solution for the **producer** is to either **go to sleep** or discard data if the buffer is full.
- The next time the consumer removes an item from the buffer, it notifies the producer, who starts to fill the buffer again.

## Solution: produce and consume problem

- The solution for the **producer** is to either **go to sleep** or discard data if the buffer is full.
- The next time the consumer removes an item from the buffer, it notifies the producer, who starts to fill the buffer again.
- In the same way, the **consumer can go to sleep** if it finds the buffer to be empty. The next time the producer puts data into the buffer, it wakes up the sleeping consumer.

## Solution: produce and consume problem

- The solution for the **producer** is to either **go to sleep** or discard data if the buffer is full.
- The next time the consumer removes an item from the buffer, it notifies the producer, who starts to fill the buffer again.
- In the same way, the **consumer can go to sleep** if it finds the buffer to be empty. The next time the producer puts data into the buffer, it wakes up the sleeping consumer.
- The **solution** can be reached by means of **inter-process communication**, typically using **semaphores**.

## Solution: produce and consume problem

- The solution for the **producer** is to either **go to sleep** or discard data if the buffer is full.
- The next time the consumer removes an item from the buffer, it notifies the producer, who starts to fill the buffer again.
- In the same way, the **consumer can go to sleep** if it finds the buffer to be empty. The next time the producer puts data into the buffer, it wakes up the sleeping consumer.
- The **solution** can be reached by means of **inter-process communication**, typically using **semaphores**.
- An **inadequate solution** could result in a **deadlock** where both processes are waiting to be awakened.

# Solution: produce and consume problem

- The solution for the **producer** is to either **go to sleep** or discard data if the buffer is full.
- The next time the consumer removes an item from the buffer, it notifies the producer, who starts to fill the buffer again.
- In the same way, the **consumer can go to sleep** if it finds the buffer to be empty. The next time the producer puts data into the buffer, it wakes up the sleeping consumer.
- The **solution** can be reached by means of **inter-process communication**, typically using **semaphores**.
- An **inadequate solution** could result in a **deadlock** where both processes are waiting to be awakened.
- The problem can also be **generalized** to have **multiple producers and consumers**.

# Channel in Julia

- A **channel** can be visualized as a **pipe**, i.e., it has a **write end** and **read end**.

# Channel in Julia

- A **channel** can be visualized as a **pipe**, i.e., it has a **write end** and **read end**.
- **Multiple writers** in different tasks **can write** to the same channel **concurrently** via **put!** calls.

# Channel in Julia

- A **channel** can be visualized as a **pipe**, i.e., it has a **write end** and **read end**.
- **Multiple writers** in different tasks **can write** to the same channel **concurrently** via **put!** calls.
- **Multiple readers** in different tasks can **read data concurrently** via **take!** calls.

# Channel operations 1/2

- Channels are created via the `Channel{T}(sz)` constructor. The channel will only hold objects of type `T`.

# Channel operations 1/2

- Channels are created via the `Channel{T}(sz)` constructor. The channel will only hold objects of type `T`.
  - If the `type is not specified`, the channel can hold objects of `any type`.

# Channel operations 1/2

- Channels are created via the `Channel{T}(sz)` constructor. The channel will only hold objects of type T.
  - If the type is not specified, the channel can hold objects of any type.
  - sz refers to the maximum number of elements that can be held in the channel at any time.

# Channel operations 1/2

- Channels are created via the `Channel{T}(sz)` constructor. The channel will only hold objects of type T.
  - If the type is not specified, the channel can hold objects of any type.
  - sz refers to the maximum number of elements that can be held in the channel at any time.
  - For example, `Channel(32)` creates a channel that can hold a maximum of 32 objects of any type.

# Channel operations 1/2

- Channels are created via the `Channel{T}(sz)` constructor. The channel will only hold objects of type T.
  - If the type is not specified, the channel can hold objects of any type.
  - sz refers to the maximum number of elements that can be held in the channel at any time.
  - For example, `Channel(32)` creates a channel that can hold a maximum of 32 objects of any type.
  - A `Channel{MyType}(64)` can hold up to 64 objects of MyType at any time.

## Channel operations 2/2

- If a Channel is empty, readers (on a `take!` call) will block until data is available.

## Channel operations 2/2

- If a Channel is empty, readers (on a `take!` call) will block until data is available.
- If a Channel is full, writers (on a `put!` call) will block until space becomes available.

## Channel operations 2/2

- If a Channel is empty, readers (on a `take!` call) will block until data is available.
- If a Channel is full, writers (on a `put!` call) will block until space becomes available.
- `isready` tests for the presence of any object in the channel, while `wait` waits for an object to become available.

## Channel operations 2/2

- If a Channel is empty, readers (on a `take!` call) will block until data is available.
- If a Channel is full, writers (on a `put!` call) will block until space becomes available.
- `isready` tests for the presence of any object in the channel, while `wait` waits for an object to become available.
- A Channel is in an open state initially. This means that it can be read from and written to freely via `take!` and `put!` calls.

## Channel operations 2/2

- If a Channel is empty, readers (on a `take!` call) will block until data is available.
- If a Channel is full, writers (on a `put!` call) will block until space becomes available.
- `isready` tests for the presence of any object in the channel, while `wait` waits for an object to become available.
- A Channel is in an open state initially. This means that it can be read from and written to freely via `take!` and `put!` calls.
- `close` closes a Channel. On a closed Channel, `put!` will fail.

## Example from Julia docs

- Consider a simple example using channels for inter-task communication.

## Example from Julia docs

- Consider a simple example using channels for inter-task communication.
- We start 4 tasks to process data from a single jobs channel.

## Example from Julia docs

- Consider a simple example using channels for inter-task communication.
- We start 4 tasks to process data from a single jobs channel.
- Jobs, identified by an id (job\_id), are written to the channel.

## Example from Julia docs

- Consider a simple example using channels for inter-task communication.
- We start 4 tasks to process data from a single jobs channel.
- Jobs, identified by an id (job\_id), are written to the channel.
- Each task in this simulation reads a job\_id, waits for a random amount of time and writes back a tuple of job\_id and the simulated time to the results channel.

## Example from Julia docs

- Consider a simple example using channels for inter-task communication.
- We start 4 tasks to process data from a single jobs channel.
- Jobs, identified by an id (job\_id), are written to the channel.
- Each task in this simulation reads a job\_id, waits for a random amount of time and writes back a tuple of job\_id and the simulated time to the results channel.
- Finally all the results are printed out.

## Example from Julia docs

- Consider a simple example using channels for inter-task communication.
- We start 4 tasks to process data from a single jobs channel.
- Jobs, identified by an id (job\_id), are written to the channel.
- Each task in this simulation reads a job\_id, waits for a random amount of time and writes back a tuple of job\_id and the simulated time to the results channel.
- Finally all the results are printed out.

## Example from Julia docs

- Consider a simple example using channels for inter-task communication.
- We start 4 tasks to process data from a single jobs channel.
- Jobs, identified by an id (job\_id), are written to the channel.
- Each task in this simulation reads a job\_id, waits for a random amount of time and writes back a tuple of job\_id and the simulated time to the results channel.
- Finally all the results are printed out.

```
julia> const jobs = Channel{Int}(32);
```

```
julia> const results = Channel{Tuple}(32);
```

```
julia> function do_work()
    for job_id in jobs
        exec_time = rand()
        sleep(exec_time) # simulates elapsed time doing actual work
                           # typically performed externally.
        put!(results, (job_id, exec_time))
    end
end;
```

## Example from Julia docs

- Consider a simple example using channels for inter-task communication.

## Example from Julia docs

- Consider a simple example using channels for inter-task communication.
- We start 4 tasks to process data from a single jobs channel.

## Example from Julia docs

- Consider a simple example using channels for inter-task communication.
- We start 4 tasks to process data from a single jobs channel.
- Jobs, identified by an id (job\_id), are written to the channel.

## Example from Julia docs

- Consider a simple example using channels for inter-task communication.
- We start 4 tasks to process data from a single jobs channel.
- Jobs, identified by an id (job\_id), are written to the channel.
- Each task in this simulation reads a job\_id, waits for a random amount of time and writes back a tuple of job\_id and the simulated time to the results channel.

## Example from Julia docs

- Consider a simple example using channels for inter-task communication.
- We start 4 tasks to process data from a single jobs channel.
- Jobs, identified by an id (job\_id), are written to the channel.
- Each task in this simulation reads a job\_id, waits for a random amount of time and writes back a tuple of job\_id and the simulated time to the results channel.
- Finally all the results are printed out.

## Example from Julia docs

- Consider a simple example using channels for inter-task communication.
- We start 4 tasks to process data from a single jobs channel.
- Jobs, identified by an id (job\_id), are written to the channel.
- Each task in this simulation reads a job\_id, waits for a random amount of time and writes back a tuple of job\_id and the simulated time to the results channel.
- Finally all the results are printed out.

## Example from Julia docs

- Consider a simple example using channels for inter-task communication.
- We start 4 tasks to process data from a single jobs channel.
- Jobs, identified by an id (job\_id), are written to the channel.
- Each task in this simulation reads a job\_id, waits for a random amount of time and writes back a tuple of job\_id and the simulated time to the results channel.
- Finally all the results are printed out.

```
julia> n = 12;

julia> @schedule make_jobs(n); # feed the jobs channel with "n" jobs

julia> for i in 1:4 # start 4 tasks to process requests in parallel
       @schedule do_work()
    end

julia> @elapsed while n > 0 # print out results
           job_id, exec_time = take!(results)
           println("$job_id finished in $(round(exec_time,2)) seconds")
           n = n - 1
       end
```

# Parallel implementation: Spatial\_arrangement in LARLIB.jl

## define channels

```

function frag_face_channel(in_chan, out_chan, V, EV, FE, sp_idx)
    run_loop = true
    while run_loop
        sigma = take!(in_chan)
        if sigma != -1
            put!(out_chan, frag_face(V, EV, FE, sp_idx, sigma))
        else
            run_loop = false
        end
    end
end

function frag_face(V, EV, FE, sp_idx, sigma)
    vs_num = size(V, 1)

    sigmavs = (abs.(FE[sigma:sigma,:])*abs.(EV))[1,:].nzind
    sV = V[sigmavs, :]
    sEV = EV[FE[sigma, :].nzind, sigmavs]

    M = submanifold_mapping(sV)
    tV = (V_ones(vs_num)\*M)\cdot 1.31

```

## spatial\_arrangement

```
function spatial_arrangement(V::Verts, EV::Cells, FE::Cells; multiproc=false)
    fs_num = size(FE, 1)
    sp_idx = spatial_index(V, EV, FE)

    rV = Verts(0,3)
    rEV = spzeros(Int8,0,0)
    rFE = spzeros(Int8,0,0)

    if (multiproc == true)
        in_chan = RemoteChannel(()->Channel{Int64}(0))
        out_chan = RemoteChannel(()->Channel{Tuple}{}(0))

        @schedule begin
            for sigma in 1:fs_num
                put!(in_chan, sigma)
            end
            for p in workers()
                put!(in_chan, -1)
            end
        end
    end
end
```

## spatial\_arrangement

```
else
    for sigma in 1:fs_num
        print(sigma, "/", fs_num, "\r")
        nV, nEV, nFE = frag_face(V, EV, FE, sp_idx, sigma)
        rV, rEV, rFE = skel_merge(rV, rEV, rFE, nV, nEV, nFE)
    end
end

rV, rEV, rFE = merge_vertices(rV, rEV, rFE)
```

## Look at the code

[https://github.com/cvdlab/LARLIB.jl/blob/master/src/spatial\\_arrangement.jl](https://github.com/cvdlab/LARLIB.jl/blob/master/src/spatial_arrangement.jl)